

Insurance companies invest a lot of time and money into optimizing their pricing and accurately estimating the likelihood that customers will make a claim. In many countries insurance it is a legal requirement to have car insurance in order to drive a vehicle on public roads, so the market is very large!

(Source: https://www.accenture.com/_acnmedia/pdf-84/accenture-machine-learning-insurance.pdf)

Knowing all of this, On the Road car insurance have requested your services in building a model to predict whether a customer will make a claim on their insurance during the policy period. As they have very little expertise and infrastructure for deploying and monitoring machine learning models, they've asked you to identify the single feature that results in the best performing model, as measured by accuracy, so they can start with a simple model in production.

They have supplied you with their customer data as a csv file called `car_insurance.csv`, along with a table detailing the column names and descriptions below.

The dataset

Column	Description
<code>id</code>	Unique client identifier
<code>age</code>	Client's age: <code>0</code> : 16-25 <code>1</code> : 26-39 <code>2</code> : 40-64 <code>3</code> : 65+
<code>gender</code>	Client's gender: <code>0</code> : Female <code>1</code> : Male
<code>driving_experience</code>	Years the client has been driving: <code>0</code> : 0-9 <code>1</code> : 10-19 <code>2</code> : 20-29 <code>3</code> : 30+
<code>education</code>	Client's level of education: <code>0</code> : No education <code>1</code> : High school <code>2</code> : University
<code>income</code>	Client's income level: <code>0</code> : Poverty <code>1</code> : Working class <code>2</code> : Middle class <code>3</code> : Upper class
<code>credit_score</code>	Client's credit score (between zero and one)
<code>vehicle_ownership</code>	Client's vehicle ownership status: <code>0</code> : Does not own their vehicle (paying off finance) <code>1</code> : Owns their vehicle
<code>vehcile_year</code>	Year of vehicle registration: <code>0</code> : Before 2015 <code>1</code> : 2015 or later
<code>married</code>	Client's marital status: <code>0</code> : Not married <code>1</code> : Married
<code>children</code>	Client's number of children
<code>postal_code</code>	Client's postal code
<code>annual_mileage</code>	Number of miles driven by the client each year
<code>vehicle_type</code>	Type of car: <code>0</code> : Sedan <code>1</code> : Sports car
<code>speeding_violations</code>	Total number of speeding violations received by the client
<code>duis</code>	Number of times the client has been caught driving under the influence of alcohol
<code>past_accidents</code>	Total number of previous accidents the client has been involved in

Column	Description
<code>outcome</code>	Whether the client made a claim on their car insurance (response variable): 0: No claim 1: Made a claim

Random Forest Classification — Required Steps

- Import the required libraries**
 - Pandas
 - Visualization libraries
 - Scikit-learn preprocessing, pipeline, Random Forest, and evaluation tools
- Load the `car_insurance.csv` dataset**
 - Display the first few rows.
- Explore the dataset**
 - Check `shape`, `info()`, and `describe()`.
 - Check for missing values.
 - Identify numerical and categorical columns.
 - Visualize the `outcome` target distribution.
- Prepare features and target**
 - Create `X` containing the input features.
 - Create `y` containing `outcome`.
 - Decide whether the `id` column should be removed.
- Identify feature types**
 - Numerical
 - Categorical
 - Ordinal, where appropriate
- Preprocess the features**
 - Apply suitable encoding/transformation.
 - Organize preprocessing using `ColumnTransformer`.
- Split the data**
 - 80% training data.
 - 20% testing data.
 - Use `random_state`.
- Create a machine learning Pipeline**
 - Combine the `ColumnTransformer` preprocessing with `RandomForestClassifier`.
- Train the Random Forest**
 - Fit the pipeline using `X_train` and `y_train`.
- Make predictions**

- Predict `X_test`.
- Store the result as `y_pred`.

11. Evaluate the model

- Accuracy
- Confusion Matrix
- Precision
- Recall
- F1-score
- Classification Report

12. Write a conclusion

- Explain the preprocessing performed.
- State the model's performance.
- Interpret the evaluation metrics.
- Conclude how well the Random Forest predicts insurance claims.

Overall workflow:

Explore → Prepare X/y → Preprocess → Split → Pipeline → Random Forest
→ Train → Predict → Evaluate → Conclusion

1

```
# Import required modules
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.pipeline import Pipeline
from sklearn.impute import SimpleImputer
from sklearn.compose import ColumnTransformer
from sklearn.preprocessing import OneHotEncoder, StandardScaler,
RobustScaler, OrdinalEncoder

# ML libraries
from sklearn.tree import DecisionTreeClassifier,
DecisionTreeRegressor, plot_tree
from sklearn.model_selection import train_test_split, cross_val_score,
GridSearchCV, RandomizedSearchCV
from sklearn.metrics import
accuracy_score, ConfusionMatrixDisplay, classification_report, mean_squar
ed_error, mean_absolute_error, r2_score, confusion_matrix, accuracy_score,
precision_score, recall_score, f1_score
from sklearn.ensemble import
RandomForestClassifier, RandomForestRegressor
from sklearn.metrics import roc_curve, auc
from scipy.stats import randint
# Start coding!
```

2

```
df = pd.read_csv('car_insurance.csv')
```

3

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 10000 entries, 0 to 9999
```

```
Data columns (total 18 columns):
```

#	Column	Non-Null Count	Dtype
0	id	10000 non-null	int64
1	age	10000 non-null	int64
2	gender	10000 non-null	int64
3	driving_experience	10000 non-null	object
4	education	10000 non-null	object
5	income	10000 non-null	object
6	credit_score	9018 non-null	float64
7	vehicle_ownership	10000 non-null	float64
8	vehicle_year	10000 non-null	object
9	married	10000 non-null	float64
10	children	10000 non-null	float64
11	postal_code	10000 non-null	int64
12	annual_mileage	9043 non-null	float64
13	vehicle_type	10000 non-null	object
14	speeding_violations	10000 non-null	int64
15	duis	10000 non-null	int64
16	past_accidents	10000 non-null	int64
17	outcome	10000 non-null	float64

```
dtypes: float64(6), int64(7), object(5)
```

```
memory usage: 1.4+ MB
```

```
df.describe()
```

	id	age	gender	credit_score	\
count	10000.000000	10000.000000	10000.000000	9018.000000	
mean	500521.906800	1.489500	0.499000	0.515813	
std	290030.768758	1.025278	0.500024	0.137688	
min	101.000000	0.000000	0.000000	0.053358	
25%	249638.500000	1.000000	0.000000	0.417191	
50%	501777.000000	1.000000	0.000000	0.525033	
75%	753974.500000	2.000000	1.000000	0.618312	
max	999976.000000	3.000000	1.000000	0.960819	

	vehicle_ownership	married	children	postal_code	\
count	10000.000000	10000.000000	10000.000000	10000.000000	
mean	0.697000	0.498200	0.688800	19864.548400	
std	0.459578	0.500022	0.463008	18915.613855	
min	0.000000	0.000000	0.000000	10238.000000	

25%	0.000000	0.000000	0.000000	10238.000000
50%	1.000000	0.000000	1.000000	10238.000000
75%	1.000000	1.000000	1.000000	32765.000000
max	1.000000	1.000000	1.000000	92101.000000


	annual_mileage	speeding_violations	duis
past_accidents \			
count	9043.000000	10000.000000	10000.000000
mean	11697.003207	1.482900	0.23920
std	2818.434528	2.241966	0.55499
min	2000.000000	0.000000	0.000000
25%	10000.000000	0.000000	0.000000
50%	12000.000000	0.000000	0.000000
75%	14000.000000	2.000000	0.000000
max	22000.000000	22.000000	6.000000

	outcome
count	10000.000000
mean	0.313300
std	0.463858
min	0.000000
25%	0.000000
50%	0.000000
75%	1.000000
max	1.000000

df.shape

(10000, 18)

```
sns.countplot(x='outcome', data=df)
plt.title('Target Outcome Distribution')
plt.show()
```



4

```
X = df.drop(columns=['id','outcome'])  
y = df['outcome']
```

5

```
num_cols = ['credit_score', 'annual_mileage', 'speeding_violations',  
            'duis', 'past_accidents']  
cat_cols = ['age', 'gender', 'driving_experience', 'education',  
            'income',  
            'vehicle_ownership', 'vehicle_year', 'married',  
            'children',  
            'postal_code', 'vehicle_type']
```

6

```
num_trans = Pipeline(steps=[  
    ('imputer', SimpleImputer(strategy='median')),  
    ('scaler', RobustScaler())  
])
```

```
cat_trans = Pipeline(steps=[
    ('imputer', SimpleImputer(strategy='constant', fill_value='missing')),
    ('encoder', OneHotEncoder())
])

preprocessor = ColumnTransformer(transformers=[
    ('num', num_trans, num_cols),
    ('cat', cat_trans, cat_cols)
])
```

7

```
X_train, X_test, y_train, y_test =
train_test_split(X, y, test_size=0.2, random_state=67)
```

8

```
rf_clf = Pipeline(steps=[
    ('preprocessor', preprocessor),

    ('classifier', RandomForestClassifier(n_estimators=100, random_state=67)
    ])

param_dist = {
    'classifier__max_depth': [3, 5, 7, 10, None],
    'classifier__min_samples_split': [2, 5, 10],
    'classifier__min_samples_leaf': randint(1, 10),
    'classifier__criterion': ['gini', 'entropy', 'log_loss'],
    'classifier__n_estimators': randint(50, 200)
}

random_search = RandomizedSearchCV(
    estimator=rf_clf,
    param_distributions=param_dist,
    cv=5,
    scoring='accuracy',
    random_state=42,
    n_jobs=-1,
    verbose=1
)

random_search.fit(X_train, y_train)

Fitting 5 folds for each of 10 candidates, totalling 50 fits

RandomizedSearchCV(cv=5,
                    estimator=Pipeline(steps=[('preprocessor',
```

```

ColumnTransformer(transformers=[('num',
Pipeline(steps=[('imputer',
SimpleImputer(strategy='median')),
('scaler',
RobustScaler())]),
['credit_score',
'annual_mileage',
'speeding_violations',
'duis',
'past_accidents']),
('cat',
Pipeline(steps=[('imputer',
SimpleImputer(fill_value='missing',
strate...
'log_loss'],
'classifier__max_depth': [3,
5, 7, 10,
None],
'classifier__min_samples_leaf':
<scipy.stats._distn_infrastructure.rv_discrete_frozen object at
0x0000027A36C07050>,
'classifier__min_samples_split': [2, 5,
10],
'classifier__n_estimators':
<scipy.stats._distn_infrastructure.rv_discrete_frozen object at
0x0000027A39844150>}},
random_state=42, scoring='accuracy', verbose=1)

```

9 and 10

```

rf_clf.fit(X_train, y_train)
y_pred = rf_clf.predict(X_test)

```


11

```
acc = accuracy_score(y_test, y_pred)
prec = precision_score(y_test, y_pred)
recall = recall_score(y_test, y_pred)
fs = f1_score(y_test, y_pred)
```


```
print(f"Accuracy score: {acc:.4f}")
print(f"Precision score: {prec:.4f}")
print(f"Recall score: {recall:.4f}")
print(f"F1 score: {fs:.4f}")
```

Accuracy score: 0.8420
Precision score: 0.7635
Recall score: 0.7324
F1 score: 0.7476

```
print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0.0	0.88	0.89	0.89	1361
1.0	0.76	0.73	0.75	639
accuracy			0.84	2000
macro avg	0.82	0.81	0.82	2000
weighted avg	0.84	0.84	0.84	2000

```
cm = confusion_matrix(y_test, y_pred)
display =
ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=['No Claim
(0)', 'Claim (1)'])
display.plot()
plt.title('RF Confusion Matrix')
plt.show()
```



```
df['outcome'].value_counts()
```


```
outcome
0.0    6867
1.0    3133
Name: count, dtype: int64
```

ROC and AOC Visualisation

```
y_prob = rf_clf.predict_proba(X_test)[:,-1]
fpr,tpr,thresholds = roc_curve(y_test,y_prob)

roc_auc = auc(fpr,tpr)

plt.figure()
plt.plot(fpr,tpr,label='ROC curve (area = %0.2f)'%)
plt.plot([0,1],[0,1],linestyle='--')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC (Receiver Operating Characteristics)')
plt.legend()
plt.show()
```



12

Explain the preprocessing performed.

The missing numerical features were imputed with mean using `StandardScaler`, while the categorical features were imputed with mode and encoded using `OneHotEncoder`. All preprocessing steps were done in a `ColumnTransformer` inside a pipeline to prevent data leakage.

State the model's performance.

The `RandomForestClassifier` achieved an accuracy score of 0.8365, outperforming the baseline single-feature models (`driving_experience`)

Interpret the evaluation metrics.

Class 0 achieved 0.87 precision and F1-score of 0.88. The model is highly accurate at identifying low-risk policyholders who will not file a claim.

Class 1 achieved 0.76 precision and F1-score of 0.74. When the model flags a customer as likely to make a claim, it is correct 76% of the time and successfully identifies 72% of all actual claims filled.

Conclude how well the RF predicts claims

The RF model demonstrates strong predictive capability for evaluating insurance risk, by combining demographic, historical and vehicle features into an ensemble model. The company can accurately forecast claim probabilities and optimize risk assessments using the model above, rather than relying on a single attribute.